

Co-expression analysis of RNA-seq data with the HTSCluster package

Andrea Rau¹, Cathy Maugis-Rabousseau, Marie-Laure Martin-Magniette, and Gilles Celeux

¹andrea.rau@jouy.inra.fr

HTSCluster version 2.0.4

Abstract

This vignette illustrates the use of the *HTSCluster* package through a worked example of a co-expression analysis using the human RNA-seq data from Sultan *et al.* (2008) [1]. For a full presentation of the statistical method, please see our paper [2].

Contents

1	Input data	1
2	Identifying co-expressed genes	2
2.1	Model description	2
2.1.1	Poisson mixture model	2
2.1.2	Inference	3
2.1.3	Model selection	3
2.2	Co-expression analysis of Sultan <i>et al.</i> (2008) data	3
2.2.1	Model selection	4
2.2.2	Visualizing results	5
3	Further reading	7
4	Session Info	7

1 Input data

In this vignette, we will work with the gene-level read counts from the Sultan *et al.* (2008) data [1], which may be found in the *HTSFilter* Bioconductor package [3]. These data were obtained from a human embryonic kidney (HEK293T) and a Ramos B cell line, with two biological replicates in each experimental condition. The raw read counts for 9010 genes and phenotype tables were originally obtained from the ReCount online resource [4].

We begin by loading the necessary packages, data, and phenotypic information for the analysis.

```
> library(HTSCluster)
> library(HTSFilter)
> library(Biobase)
> data(sultan)
> conds <- as.vector(phenoData(sultan)$cell.line)
> y <- exprs(sultan)
```

As an additional pre-processing step, we apply the data-based filter proposed in the *HTSFilter* package [3] to remove weakly expressed genes across the two conditions. This approach identifies a filtering threshold by maximizing a global Jaccard similarity index calculated between replicates within each condition. After applying this threshold (see Figure 1), 4956 genes were retained for the subsequent co-expression analysis.

```
> y.filter <- HTSFilter(y, conds, norm="TMM")
> table(y.filter$on) ## 4054 off, 4956 on
```




Figure 1: Global Jaccard similarity index calculated over various filtering thresholds for normalized counts for the [1] data via the HTSFilter package [3]. The data-driven filtering threshold in this case is equal to $s^* = 12.02$.

```
0 1
4054 4956
```

```
> dat.select <- y.filter$filterData
```

2 Identifying co-expressed genes

2.1 Model description

The following description closely follows that provided in our main paper [2].

Let Y_{ijl} be the random variable corresponding to the digital gene expression measure (DGE) for biological entity i ($i = 1, \dots, n$) of condition j ($j = 1, \dots, d$) in biological replicate l ($l = 1, \dots, r_j$), with y_{ijl} being the corresponding observed value of Y_{ijl} . Let $q = \sum_{j=1}^d r_j$ be the total number of variables (all replicates in all conditions) in the data, such that $\mathbf{y} = (y_{ijl})$ is the $n \times q$ matrix of the DGE for all observations and variables, and \mathbf{y}_i is the q -dimensional vector of DGE for all variables of observation i . We use dot notation to indicate summations in various directions, e.g., $y_{\cdot j l} = \sum_i y_{ijl}$, $y_{i \cdot} = \sum_j \sum_l y_{ijl}$, and so on.

2.1.1 Poisson mixture model

To cluster RNA-seq data, we consider a model-based clustering procedure based on mixture of Poisson distributions. The data \mathbf{y} are assumed to come from K distinct subpopulations (clusters), each of which is modeled separately:

$$f(\mathbf{y}; K, \Psi_K) = \prod_{i=1}^n \sum_{k=1}^K \pi_k f_k(\mathbf{y}_i; \theta_{ik})$$

where $\Psi_K = (\pi_1, \dots, \pi_{K-1}, \theta')'$, θ' contains all of the parameters in $\{\theta_{ik}\}_{i,k}$ and $\pi = (\pi_1, \dots, \pi_K)'$ are the mixing proportions, with $\pi_k \in (0, 1)$ for all k and $\sum_{k=1}^K \pi_k = 1$. Samples are assumed to be independent conditionally on the components:

$$f_k(\mathbf{y}_i; \theta_{ik}) = \prod_{j=1}^d \prod_{l=1}^{r_j} \mathcal{P}(y_{ijl}; \mu_{ijlk}),$$

where $\mathcal{P}(\cdot)$ denotes the standard Poisson probability mass function and $\theta_{ik} = \{\mu_{ijlk}\}_{j,l}$.

Each mean μ_{ijlk} is parameterized by

$$\mu_{ijlk} = w_i s_{jl} \lambda_{jk}$$

where $w_i = y_{i..}$ corresponds to the overall expression level of observation i (e.g., weakly to strongly expressed) and s_{jl} represents the normalized library size for replicate l of condition j , such that $\sum_{j,l} s_{jl} = 1$. These normalization factors take into account the fact that the number of reads expected to map to a particular gene depends not only on its expression level, but also on the library size (overall number of mapped reads) and the overall composition of the RNA population being sampled. We note that $\{s_{jl}\}_{j,l}$ are estimated from the data prior to fitting the model, and like the overall expression levels w_i , they are subsequently considered to be fixed in the Poisson mixture model. Finally, the unknown parameter vector $\lambda_k = (\lambda_{1k}, \dots, \lambda_{dk})$ corresponds to the clustering parameters that define the profiles of the genes in cluster k across all biological conditions.

2.1.2 Inference

To estimate mixture parameters $\Psi_K = (\pi, \lambda_1, \dots, \lambda_K)$ by computing the maximum likelihood estimate (MLE), an Expectation-Maximization (EM) algorithm is considered. After initializing the parameters $\Psi_K^{(0)}$ and $\mathbf{z}^{(0)}$ by a so-called Small-EM strategy, the E-step at iteration b corresponds to computing the conditional probability that an observation i arises from the k th component for the current value of the mixture parameters:

$$t_{ik}^{(b)} = \frac{\pi_k^{(b)} f_k(\mathbf{y}_i; \boldsymbol{\theta}_{ik}^{(b)})}{\sum_{m=1}^K \pi_m^{(b)} f_m(\mathbf{y}_i; \boldsymbol{\theta}_{im}^{(b)})}$$

where $\boldsymbol{\theta}_{ik}^{(b)} = \{w_i s_{jl} \lambda_{jk}^{(b)}\}_{j,l}$. Then, in the M-step the mixture parameter estimates are updated to maximize the expected value of the completed likelihood, which leads to weighting the observation i for group k with the conditional probability $t_{ik}^{(b)}$. Thus,

$$\pi_k^{(b+1)} = \frac{1}{n} \sum_{i=1}^n t_{ik}^{(b)} \quad \text{and} \quad \lambda_{jk}^{(b+1)} = \frac{\sum_{i=1}^n t_{ik}^{(b)} y_{ij.}}{s_{j.} \sum_{i=1}^n t_{ik}^{(b)} y_{i..}},$$

since $w_i = y_{i..}$. Note that at each iteration of the EM algorithm, we obtain that $\sum_{j=1}^d \lambda_{jk}^{(b)} s_{j.} = 1$. Thus $\lambda_{jk}^{(b)} s_{j.}$ can be interpreted as the proportion of reads that are attributed to condition j in cluster k , after accounting for differences due to library size; this proportion is shared among the replicates of condition j according to their respective library sizes s_{jl} .

2.1.3 Model selection

For model selection (i.e., the choice of the number of clusters K), we make use of the so-called *slope heuristics*, which is a data-driven method to calibrate a penalized criterion whose penalties are known only up to a multiplicative factor. In practice, the slope heuristics approach involves determining the minimal penalty given the number of free parameters; this is typically performed using either the *dimension jump* (Djump), which calibrates the penalty according to the greatest jump in complexity as a function of the multiplicative constant in the penalty, or the *data-driven slope estimation* (DDSE), which is based on the expectation of a linear relationship between the penalty shape and likelihood for the most complex models. Both of these approaches are implemented in the *capushe* R package, available on CRAN [5].

After calibrating the penalty using the slope heuristics, the number of selected clusters \hat{K} then corresponds to the value of K minimizing the associated penalized criterion. Based on $\hat{\Psi}_{\hat{K}}$, each observation i is assigned to the component maximizing the conditional probability \hat{t}_{ik} (i.e., using the so-called MAP rule).

2.2 Co-expression analysis of Sultan *et al.* (2008) data

In the interest of computational time, we perform a single run of HTScluster for $K = 10, \dots, 20$ clusters, using the Trimmed Means of M-values (TMM) normalization [6], the splitting-small EM strategy described in the main paper and a convergence cutoff (in terms of the difference in loglikelihoods from one iteration to the next) equal to 10^{-4} ; we note that this cutoff is fixed to be larger than the default (10^{-6}) simply to speed up calculations for this vignette. Because the settings here differ from those used in the full analysis (i.e., higher cutoff and smaller set of models), the results presented here differ slightly from those presented in the main paper [2].

```
> set.seed(12345)
> PMM <- PoisMixClusWrapper(y=dat.select, gmin=10, gmax=20,
+   conds=conds, split.init=TRUE, cutoff=10e-4, lib.type="TMM")
```

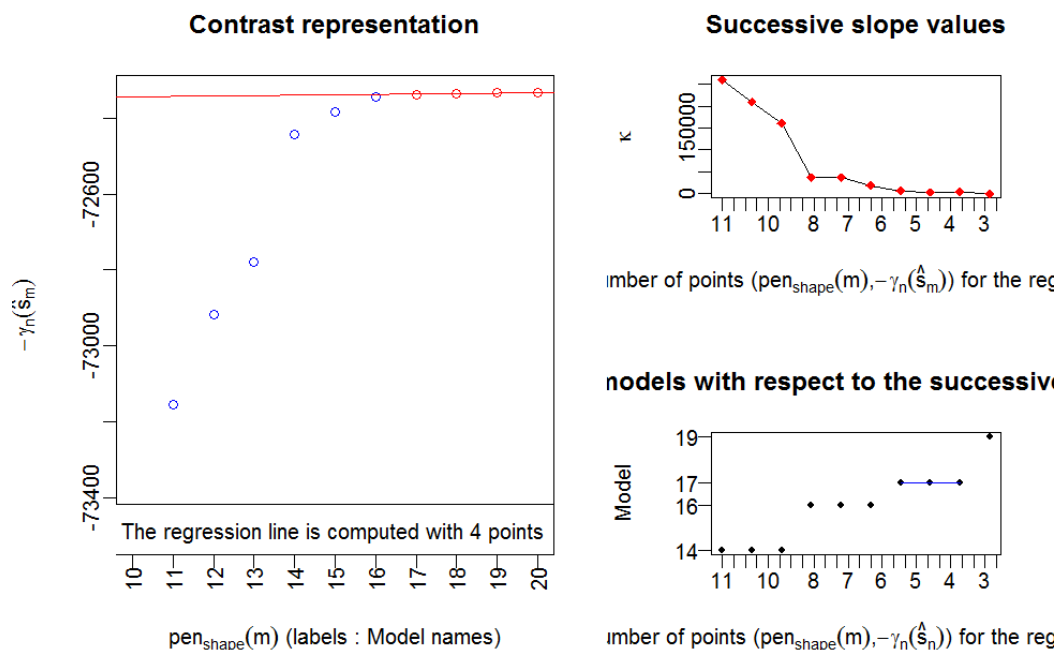


Figure 2: Diagnostic plots provided by *capushe* package for the DDSE approach; see [5] for additional information.

```
Running g = 10 ...
Running g = 11 ...
Running g = 12 ...
Running g = 13 ...
Running g = 14 ...
Running g = 15 ...
Running g = 16 ...
Running g = 17 ...
Running g = 18 ...
Running g = 19 ...
Running g = 20 ...
```

Note that a warning is produced by *capushe* if the models returned by the DDSE and Djump approaches are not the same.

2.2.1 Model selection

The models selected by the DDSE and Djump criteria may be accessed as follows:

```
> mod.DDSE <- PMM$DDSE.results
> mod.Djump <- PMM$Djump.results
```

The *capushe* package provides diagnostic plots for the slope heuristics in order to ensure that sufficiently complex models have been considered. We note that in this particular example, the set of models under consideration is too small ($K = 10, \dots, 20$) to appropriately make use of the slope heuristics; however, for illustration we illustrate how to produce diagnostic plots for the slope heuristics.

Results from the *capushe* package over the set of models fit by *HTScluster* may be found in the *capushe* subset of objects of class *HTSclusterWrapper* (i.e., the output of the *PoisMixClusWrapper* function). To access the results, diagnostic plots (see Figure 2), and the selected model dimension of the DDSE method, the following code may be used.

```
> DDSE <- PMM$capushe@DDSE          ## DDSE results
> plot(DDSE, newwindow=F, ask=F)    ## DDSE diagnostic plots
> DDSE@model                        ## Model selected by DDSE
```

```
[1] "17"
```

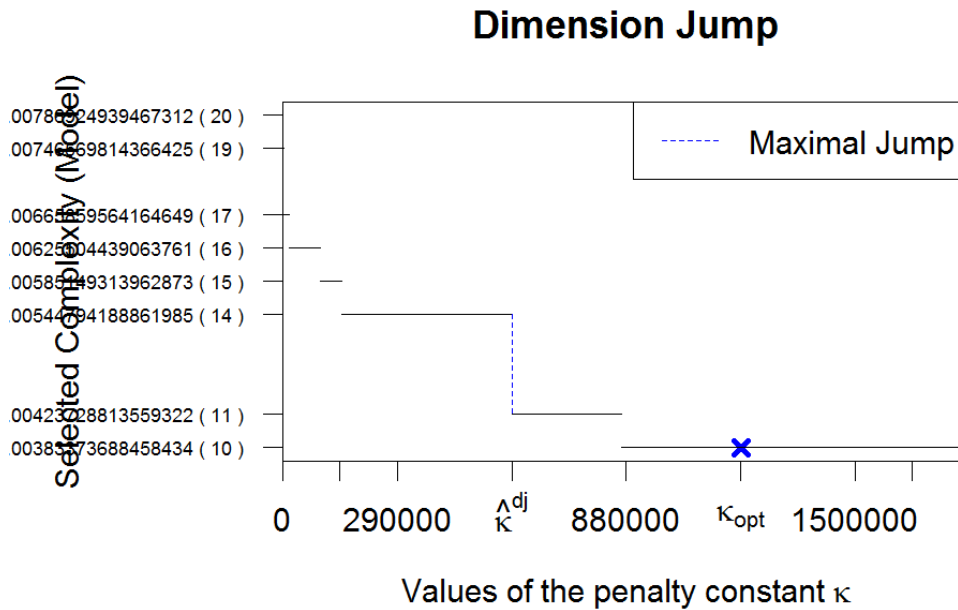


Figure 3: Diagnostic plots provided by *capushe* package for the Djump approach; see [5] for additional information.

To access the results, diagnostic plots (see Figure 3), and the selected model dimension of the Djump method, the following code may be used.

```
> Djump <- PMM$capushe@Djump      ## Djump results
> plot(Djump, newwindow=F, ask=F)  ## Djump diagnostic plots
> Djump@model                      ## Model selected by Djump

[1] "10"
```

Also, note that all *capushe* diagnostic plots may be obtained directly from the HTSclusterWrapper object using the following command:

```
> ## Not run:
> ## plot(PMM, graphs="capushe")
```

2.2.2 Visualizing results

For the following summarization and visualization, we will make use of the model selected by the Djump approach.

```
> mod <- PMM$Djump.results
```

A built-in summary command allows a text-based overview of the selected model, including the number of clusters, the Integrated Completed Likelihood (ICL) and Bayesian Information Criterion (BIC) values, the number of genes in each cluster, the number of genes with maximum conditional probabilities greater than 90%, the number of genes in each cluster with maximum conditional probabilities greater than 90%, and the estimated values of $\hat{\lambda}$ and $\hat{\pi}$.

```
> summary(mod)

*****
Number of clusters = 10
(ICL = -75397)
(BIC = -73584)
*****
Cluster sizes:
Cluster 1 Cluster 2 Cluster 3 Cluster 4 Cluster 5 Cluster 6 Cluster 7 Cluster 8 Cluster 9
156      252      439      201      254      731      600      556      764
```

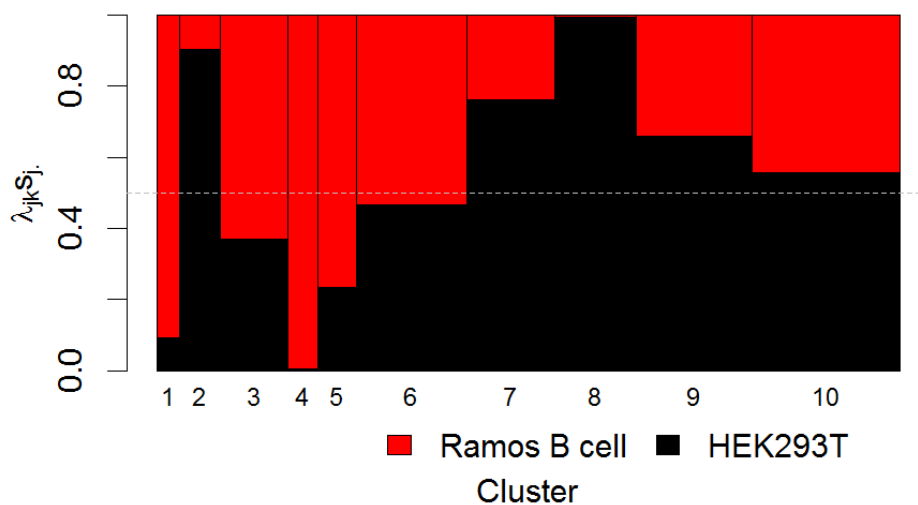


Figure 4: Visualization of overall cluster behavior for the Sultan *et al.* data. For each cluster, bar plots of $\hat{\lambda}_{jk} s_j$ are drawn for each experimental condition, where the width of each bar corresponds to the estimated proportion $\hat{\pi}_k$.

Cluster 10
1003

Number of observations with MAP > 0.90 (% of total):
2651 (53.5%)

Number of observations with MAP > 0.90 per cluster (% of total per cluster):

Cluster 1	Cluster 2	Cluster 3	Cluster 4	Cluster 5	Cluster 6	Cluster 7	Cluster 8	Cluster 9
101	166	193	187	154	276	319	495	322
(64.74%)	(65.87%)	(43.96%)	(93.03%)	(60.63%)	(37.76%)	(53.17%)	(89.03%)	(42.15%)

Cluster 10
438
(43.67%)

Lambda:

	Cluster 1	Cluster 2	Cluster 3	Cluster 4	Cluster 5	Cluster 6	Cluster 7	Cluster 8
HEK293T	0.16	1.61	0.66	0.01	0.42	0.83	1.36	1.77
Ramos B cell	2.07	0.22	1.43	2.27	1.74	1.22	0.54	0.01

	Cluster 9	Cluster 10
HEK293T	1.17	0.99
Ramos B cell	0.78	1.01

Pi:

Cluster 1	Cluster 2	Cluster 3	Cluster 4	Cluster 5	Cluster 6	Cluster 7	Cluster 8	Cluster 9
0.03	0.05	0.09	0.04	0.05	0.15	0.12	0.11	0.15

Cluster 10
0.20

The estimated values for $\hat{\lambda}$ and $\hat{\pi}$ may also be visualized using barplots, as in Figure 4, where bar widths represent the values of $\hat{\pi}$.

```
> plot(mod, graphs="lambda")
```




Figure 5: (left) Histogram of maximum conditional probabilities of cluster membership. (right) Boxplots of maximum conditional probabilities of cluster membership for the genes assigned to each cluster.

Finally, we may also examine a histogram of maximum conditional probabilities of cluster membership for all genes (Figure 5, left), as well as boxplots of maximum conditional probabilities of cluster membership for the genes assigned to each cluster (Figure 5, right). These plots help to evaluate the degree of certitude accorded by the model in assigning genes to clusters, as well as whether some clusters are attributed a greater degree of uncertainty than others. In this example, we see that most genes appear to have large maximum conditional probabilities of cluster membership (i.e., greater than 90%), and that most genes are attributed to Clusters 4 and 8 tend to be associated with large maximum conditional probabilities of cluster membership.

```
> plot(mod, graphs="map")
> plot(mod, graphs="map.bycluster")
```

Finally, the cluster labels and conditional probabilities of cluster membership assigned to each gene may be accessed using the following code:

```
> labels <- mod$labels
> probaPost <- mod$probaPost
```

3 Further reading

For additional information on the statistical method illustrated in this vignette, see [2].

4 Session Info

```
> sessionInfo()

R version 3.1.1 (2014-07-10)
Platform: x86_64-w64-mingw32/x64 (64-bit)

locale:
 [1] LC_COLLATE=French_France.1252  LC_CTYPE=French_France.1252    LC_MONETARY=French_France.1252
 [4] LC_NUMERIC=C                   LC_TIME=French_France.1252

attached base packages:
[1] parallel  stats      graphics  grDevices  utils      datasets  methods   base

other attached packages:
[1] HTSFilter_1.4.0    Biobase_2.24.0      BiocGenerics_0.10.0 HTScluster_2.0.4
[5] capushe_1.0        MASS_7.3-33

loaded via a namespace (and not attached):
 [1] annotate_1.42.1      AnnotationDbi_1.26.0 DBI_0.2-7            DESeq_1.16.0
```

[5] DESeq2_1.4.5	edgeR_3.6.7	genefilter_1.46.1	geneplotter_1.42.0
[9] GenomeInfoDb_1.0.2	GenomicRanges_1.16.4	grid_3.1.1	IRanges_1.22.10
[13] lattice_0.20-29	limma_3.20.8	locfit_1.5-9.1	plotrix_3.5-7
[17] poisson.glm.mix_1.2	RColorBrewer_1.0-5	Rcpp_0.11.2	RcppArmadillo_0.4.320.0
[21] RSQLite_0.11.4	splines_3.1.1	stats4_3.1.1	survival_2.37-7
[25] tools_3.1.1	XML_3.98-1.1	xtable_1.7-3	XVector_0.4.0

References

- [1] M. Sultan et al. A global view of gene activity and alternative splicing by deep sequencing of the human transcriptome. *Science*, 15(5891):956–60, 2008.
- [2] A. Rau et al. Co-expression analysis of high-throughput transcriptome sequencing data with Poisson mixture models. (*submitted*), 2014.
- [3] A. Rau et al. Data-based filtering for replicated high-throughput transcriptome sequencing experiments. *Bioinf.*, 29(17):2146–2152, 2013.
- [4] A. C. Frazee et al. ReCount: a multi-experiment resource of analysis-ready RNA-seq gene count datasets. *BMC Bioinformatics*, 12(449), 2011.
- [5] Jean-Patrick Baudry, Cathy Maugis, and Bertrand Michel. Slope heuristics: overview and implementation. *Stat. Comp.*, 22:455–470, 2012.
- [6] Mark D. Robinson and Alicia Oshlack. A scaling normalization method for differential expression analysis of RNA-seq data. *Genome Biology*, 11(R25), 2010.